Efficient (re)naming in Conflict-free Replicated Data types (CRDTs)

Matthieu Nicolas, Gérald Oster and Olivier Perrin

March 25, 2018

In order to design large scale distributed systems, the literature and companies increasingly adopt the optimistic replication model known as eventual consistency to replicate data among nodes. This consistency model allows replicas to temporarily diverge to be able to ensure high availability. Each node owning a copy of the data can edit it without any kind of coordination with other nodes, before propagating the changes to others. A conflict resolution mechanism is however required to handle updates generated in parallel by different replicas.

An approach which gains in popularity since a few years proposes to define Conflict-free Replicated Data types (CRDTs)[1]. These data structures behave as traditional ones, like the *Set* or the *Sequence* data structures, but are designed for a distributed usage. Their specification ensures that concurrent changes are resolved deterministically, without any kind of agreement, and that replicas eventually converge after observing all updates, thus achieving *Strong Eventual Consistency*.

In [1], the author presents two designs of CRDTs : *State-based* CRDTs and *Operations-based* CRDTs.

State-based CRDTs define data structures with monotonically increasing states with idempotent and commutative merge functions. This allows one replica to share its local updates by broadcasting its state to others. Upon the reception of the state of another replica, a node is able to update its own state by merging them, regardless of its concurrent updates. Thanks to the properties of the defined state and merge function, states can be missed or delivered multiple times : as long as the most recent state of each replica is successfully broadcast to others once, each node will converge. Thus, no assumptions are made on the network layer. However, this is achieved by broadcasting the whole state repeatedly, which may be inefficient according to the size of the data structure.

Operations-based CRDTs define data structures with a set of operations to perform updates and a partial order between these operations, usually the causal order. In addition, operations have to be designed such as concurrent operations commute. This allows to propagate local updates by broadcasting corresponding operations to other replicas. Operations are delivered according to the defined partial order and, upon delivery of an operation, a replica

updates its state by applying the received operation. In comparison to *State-based* CRDTs, this solution achieves better performances, especially regarding the bandwidth consumption. Nevertheless, it requires the network layer to keep track of the defined partial order, which may be a complex and costly task.

In both designs, to achieve convergence, CRDTs proposed in the literature mostly rely on identifiers to reference updated elements. To be globally unique, element identifiers often include the identifier of the node which generates them. But, since node identifiers grow as new nodes join the system, element identifiers grow proportionally. Furthermore, element identifiers have to comply to additional constraints according to the CRDT, for example forming a dense set. This usually results in an acceleration of their growth.

Hence, since the size of identifiers is not bounded, the size of metadata attached to each element increases over time. It thus exceeds more and more the size of data itself. This impedes the adoption of CRDTs since nodes have to broadcast and store metadata, causing the application's performances and efficiency to decrease over time.

The goal of this PhD is to address this issue by 1. proposing more efficient specifications of identifiers according to their set of constraints, 2. proposing mechanisms to rename identifiers to reduce their size.

References

[1] Marc Shapiro, Nuno Preguiça, Carlos Baquero, and Marek Zawirski. A comprehensive study of Convergent and Commutative Replicated Data Types. Research Report RR-7506, Inria – Centre Paris-Rocquencourt, January 2011.